

FuseLens: Long-Context Genomic Foundation Models for Robust DNA Breakpoint Detection

Udi Bahat
MSc Computer Science
Department of Computer Science

December 4, 2025

Abstract

Gene fusions are critical oncogenic drivers, yet identifying their precise genomic breakpoints remains computationally challenging due to repetitive regions and complex rearrangements. Traditional alignment-based methods often fail to capture long-range dependencies in DNA sequences, leading to high false-positive rates. We introduce **FuseLens**, a novel deep learning framework leveraging **HyenaDNA**, a long-convolution foundation model pretrained on the human genome. Unlike existing approaches that rely on short-read alignment or limited-context CNNs, FuseLens processes 20kb genomic windows at single-nucleotide resolution. We propose a specialized **Attention Pooling** mechanism to localize breakpoint signals and introduce a mathematically rigorous **Synthetic Negative Augmentation** strategy. Furthermore, we formalize a domain adaptation protocol to integrate real RNA-Seq data with synthetic baselines. Benchmarking on a curated dataset demonstrates state-of-the-art performance with an **AUC of 0.974** and **Recall of 91.7%**, significantly outperforming baseline expectations for sequence-based classifiers.

1 Introduction

The detection of gene fusions—hybrid genes formed by chromosomal rearrangements—is a cornerstone of precision oncology. While tools like STAR-Fusion [4] and Arriba [5] have standardized detection from RNA-Seq data, they fundamentally rely on short-read alignment. This “alignment-first” paradigm struggles with ambiguous mappings in repetitive regions (e.g., *Alu* repeats) and often fails to capture the broader genomic context (intronic sequences, regulatory motifs) that characterizes true breakpoints.

Deep learning offers an alternative “sequence-first” paradigm. Previous models like FusionAI [2] utilized Convolutional Neural Networks (CNNs) on 20kb sequences. However, standard CNNs have limited receptive fields, and Transformers scale quadratically ($O(N^2)$), making them computationally prohibitive for long genomic sequences ($> 4\text{kb}$).

We present **FuseLens**, which adapts **HyenaDNA** [3], a sub-quadratic ($O(N \log N)$) genomic foundation model, for breakpoint detection. By fine-tuning HyenaDNA’s long-range feature extraction capabilities and pairing it with a custom **Attention Pooling** head, FuseLens learns to distinguish true fusion breakpoints from biological noise with unprecedented accuracy.

2 Biological Background: Chimeric Transcripts in Cancer

Chimeric transcripts are RNA molecules formed by the joining of exons from two or more different genes. Understanding their diverse formation mechanisms is critical for distinguishing pathogenic drivers from benign artifacts. Figure ?? illustrates the four primary classes of chimeric transcript formation [1].

2.1 Mechanisms of Formation

- **Read-Through (Cis-SAGE):** This mechanism occurs at the transcriptional level when RNA polymerase fails to recognize the stop signal of the upstream gene and continues transcribing into the neighboring gene's coding region. The resulting fusion transcript is produced from an otherwise intact genomic sequence. This process is often a benign physiological event, for instance, in the case of *SLC45A3-ELK4*.
- **Chromosomal Rearrangement:** This involves permanent DNA structural alterations, where DNA breaks and rejoins incorrectly. This category includes:
 - *Intra-chromosomal* inversions (e.g., *EML4-ALK*), where a segment on a single chromosome breaks, flips, and re-ligates, joining two genes.
 - *Inter-chromosomal* translocations (e.g., *BCR-ABL1*), where segments between two different chromosomes are swapped and joined, juxtaposing the two fusion partner genes.
- **Cis-Type Fusion (Structural Variant-Driven):** These fusions occur between genes located on the same chromosome (in *cis*), but are driven by a structural variant in the DNA, such as a large deletion or duplication, that physically removes the intervening sequence and brings the two genes into direct contiguity. This results in a permanent genomic alteration that is subsequently transcribed as a single fusion transcript.

2.2 Key Oncogenic Fusions

We focus on detecting high-impact fusions characterized in the literature:

- **BCR-ABL1 (CML):** Resulting from the t(9;22) translocation (Philadelphia Chromosome). It creates a constitutively active tyrosine kinase driving Chronic Myeloid Leukemia [10].




Figure 1: **The BCR-ABL1 Fusion Event.** Reciprocal translocation t(9;22) creates a fusion gene on the derivative chromosome 22, characteristic of CML [10].

- **TMPRSS2-ERG (Prostate Cancer):** A fusion of the androgen-responsive promoter of *TMPRSS2* with the oncogene *ERG*. Found in ~50% of prostate cancers [11].
- **SLC45A3-ELK4 (Prostate Cancer):** A prominent example of **Cis-Splicing**. This chimera can form via transcriptional read-through without any underlying DNA mutation [12].

2.3 Non-Cancerous Physiological Chimeras

Not all chimeric transcripts are pathogenic. Normal physiological chimeras occur in healthy cells and can drive evolutionary adaptation or protein diversity.

- **Benign Read-Throughs:** Transcripts like *ZFAND3-ZFAND6* occur in normal human tissues via Cis-SAGe (Splicing of Adjacent Genes) to create novel protein isoforms without genomic rearrangement [15].
- **Normal Trans-Splicing:** The *JAZF1-JJAZ1* RNA chimera, originally thought to be specific to endometrial sarcoma, has been identified in normal endometrial stroma cells resulting from trans-splicing rather than DNA fusion [13].
- **Evolutionary Adaptation:** Chimerism is a mechanism of gene birth. The *Jingwei* gene in *Drosophila* originated ~2.5 million years ago from the retrotransposition of *Adh* into the *Ymp* gene, conferring new metabolic capabilities [14].

Distinguishing these benign physiological events from oncogenic drivers is a critical challenge for high-precision clinical classifiers like FuseLens.

3 Related Work

- **Alignment-Based Detectors:** Tools like STAR-Fusion [4] and Arriba [5] define the current clinical standard but are prone to high False Positive rates in complex genomic regions.
- **Deep Learning in Genomics:** Models like FusionAI [2] demonstrated the feasibility of using neural networks for structural variant detection.
- **Genomic Foundation Models:** HyenaDNA [3] and Nucleotide Transformer [7] represent a shift towards “pretraining + fine-tuning.”

4 Methods

4.1 Data Engineering: The Hard-Negative Strategy

A critical innovation of FuseLens is the rigorous definition of the “Negative” (non-chimeric) space. We define a set of canonical sequences $S = \{s_1, \dots, s_N\}$ derived from **UniProt Swiss-Prot** and apply three mathematical operators to generate “Hard Negatives.”

Let sequence s_i be a string of length L where $s_i \in \{A, T, C, G\}^L$.

1. Directionality Operator ($\mathcal{R}$):

$$\mathcal{R}(s_i) = (b_L, b_{L-1}, \dots, b_1) \quad (1)$$

Tests if the model detects biological directionality ($5' \rightarrow 3'$).

2. Composition Invariance Operator ($\mathcal{P}_{\text{rand}}$):

Let σ be a random permutation of indices $\{1, \dots, L\}$.

$$\mathcal{P}_{\text{rand}}(s_i) = (b_{\sigma(1)}, b_{\sigma(2)}, \dots, b_{\sigma(L)}) \quad (2)$$

Destroys sequence grammar while strictly preserving GC-content.

3. Artifactual Junction Operator ($\mathcal{J}$):

Simulates a random fusion between two unrelated genes s_i and s_j .

$$\mathcal{J}(s_i, s_j) = s_i[1 : \frac{L}{2}] \oplus s_j[\frac{L}{2} + 1 : L] \quad (3)$$

Teaches the model that a “junction” alone is not enough; it must be a *valid, oncogenic-like* breakpoint structure.

4.2 Model Architecture

We utilize the `hyenadna-small-32k` backbone. The core Hyena operator approximates the attention matrix via data-controlled long convolutions:

$$y = \text{Hyena}(u) = (H_1(u) * \psi) \cdot H_2(u) \quad (4)$$

Where $*$ denotes convolution, ψ is a long implicit filter, and $H(u)$ are linear projections.

Attention Pooling Head: We implement a learnable pooling layer. Let $E \in \mathbb{R}^{L \times D}$ be the output embedding, where L is the sequence length (20,480) and D is the model’s hidden embedding dimension (256). We compute a scalar attention score α_t for each nucleotide position t :

$$\alpha_t = \text{softmax}(W_a \cdot E_t + b_a) \quad (5)$$

The final sequence representation v is the weighted sum:

$$v = \sum_{t=1}^L \alpha_t E_t \quad (6)$$

4.3 Breakpoint Localization via Attention Extraction

Beyond binary classification, FuseLens is architected to perform breakpoint localization without requiring dense segmentation masks. We interpret the attention weights α_t as a probability distribution over the genomic positions $t \in [1, L]$, reflecting the model’s focus on the junction region. The predicted breakpoint coordinate $\hat{L}_{bp}$ is derived as the position of maximum attention:

$$\hat{L}_{bp} = \underset{t}{\operatorname{argmax}}(\alpha_t) \quad (7)$$

4.4 Mathematical Formulation of Data Integration

To bridge the gap between idealized synthetic training and real-world biological variability, we implement a hybrid data integration strategy. We modify the optimization objective to a **Regularized Empirical Risk Minimization (R-ERM)**:

$$\theta_{total}^* = \operatorname{argmin}_{\theta} \left[\sum_{(x,y) \in S_{syn}} \mathcal{L}(f_{\theta}(x), y) + \lambda \sum_{(x',y') \in S_{real}} \mathcal{L}(f_{\theta}(x'), y') \right] \quad (8)$$

where the second term acts as a biological constraint, forcing the decision boundary to remain valid within the noisy regime of real RNA-Seq data.

5 Benchmark Results

We evaluated FuseLens on a hold-out test set using the full augmented dataset.

The results demonstrate high robustness. Figure 2 illustrates the model’s inference characteristics on the test set.

Metric	Score
Accuracy	0.9102
Precision	0.9186
Recall	0.9174
F1 Score	0.9180
AUC	0.9744

Table 1: Final Validation Metrics (Augmented Dataset)




Figure 2: **Inference Dynamics on Test Data.** (Left) **Confidence Distribution:** The model exhibits high certainty, with the majority of positive predictions (red) clustered above 0.9 probability. (Right) **Breakpoint Localization:** The Attention Pooling mechanism successfully identifies regions of interest. The peaks indicate specific genomic indices where the model attends most strongly, correlating with putative fusion junctions.

As shown in Figure 2 (Left), the model is highly decisive, with very few samples falling in the ambiguous 0.4-0.6 probability range. The Localization plot (Right) confirms that the attention mechanism is not random; it forms distinct peaks around specific genomic coordinates (e.g., indices 10k, 14k, 18k), suggesting that FuseLens is learning to recognize specific structural motifs rather than just global sequence composition.

6 Discussion

Why Attention Pooling Matters: Traditional CNNs often struggle with the “needle in a haystack” problem. Our Attention Pooling mechanism acts as a learnable “spotlight,” allowing the model to dynamically weight the junction region higher than the intronic flanking sequences.

Synthetic Data Efficacy: Initial experiments with simple random negatives yielded high false-positive rates. The introduction of the $\mathcal{J}$ (Random Pair) and $\mathcal{P}_{\text{rand}}$ (Shuffled) operators forced the decision boundary to tighten around *structurally valid* breakpoints.

7 Limitations & Future Work

While FuseLens achieves state-of-the-art performance, we acknowledge specific constraints:

- **The Synthetic Gap:** Despite R-ERM, a distributional shift remains between synthetic “Hard Negatives” and naturally occurring biological artifacts.
- **Weakly Supervised Localization:** While FuseLens extracts breakpoint coordinates via attention pooling, these are derived from classification gradients rather than supervised regression.

7.1 Future Directions: Clinical Integration

The immediate evolution of FuseLens involves integration with **Third-Generation Sequencing** (Oxford Nanopore/PacBio). We propose a hybrid workflow where tools like **CTAT-LR-fusion** [6] serve as candidate generators, utilizing **minimap2** for initial isoform detection. FuseLens will act as a high-precision validation layer, analyzing the semantic structure of the signal data to correct alignment errors.

References

- [1] Milanesi, L. et al. “ChiTaRS 5.0: the comprehensive database of chimeric transcripts.” *Nucleic Acids Research*, 2020.
- [2] Kim, P., Tan, H., et al. “FusionAI: predicting fusion breakpoint from DNA sequence with deep learning.” *iScience*, 2021.
- [3] Nguyen, E., et al. “HyenaDNA: Long-Range Genomic Sequence Modeling at Single Nucleotide Resolution.” *NeurIPS*, 2023.
- [4] Haas, B.J., et al. “STAR-Fusion: Fast and Accurate Fusion Transcript Detection from RNA-Seq.” *bioRxiv*, 2019.
- [5] Uhrig, S., et al. “Accurate and efficient detection of gene fusions from RNA sequencing data.” *Genome Research*, 2021.
- [6] Haas, B.J., et al. “CTAT-LR-fusion: Accurate Fusion Transcript Detection from Long-Read RNA-seq.” *GitHub/Preprint*, 2024.
- [7] Dalla-Torre, H., et al. “The Nucleotide Transformer: Building and Evaluating Robust Foundation Models for Human Genomics.” *Nature Methods*, 2025.
- [8] Qin, F., et al. “Intergenically Spliced Chimeric RNAs in Cancer.” *International Journal of Cell Biology*, 2016.
- [9] Lei, Q., et al. “Evolutionary Insights into RNA trans-Splicing in Vertebrates.” *Genome Biology and Evolution*, 2016.
- [10] Ren, R. “Mechanisms of BCR-ABL in the pathogenesis of chronic myelogenous leukemia.” *Nature Reviews Cancer*, 2005.
- [11] Tomlins, S.A., et al. “Recurrent fusion of TMPRSS2 and ETS transcription factor genes in prostate cancer.” *Science*, 2005.
- [12] Zhang, Y., et al. “SLC45A3-ELK4 Chimera in Prostate Cancer: Spotlight on Cis-Splicing.” *Cancer Discovery*, 2012.
- [13] Li, H., et al. “The JAZF1-JJAZ1 gene fusion in endometrial stromal tumors: imaging and molecular analysis.” *Science*, 2008.
- [14] Long, M., et al. “The origin of new genes: glimpses from the young and old.” *Nature Reviews Genetics*, 2003.
- [15] Zhang, Y., et al. “Chimeric transcript generated by cis-splicing of adjacent genes regulates prostate cancer cell growth.” *Cancer Discovery*, 2012.